


Cloud Spanner

Setting up Google Cloud Spanner



- **Definition:** Cloud Spanner provides a highly available, SQL-compliant database solution capable of scaling horizontally on a planet-wide scale.
- **Purpose:** To Create a Cloud Spanner instance, creating a populated database, querying the data, and then exporting it to Google Cloud Storage.
- Follow the link to complete the Lab
 - Open the url – <https://console.cloud.google.com>
 - Login with your Accenture credentials (e-mail id)
 - Follow the below steps to complete your lab

Steps:

1. Open Cloud Spanner from the Google Cloud Console.
2. Go to the Create an instance page.
3. Enter an **Instance name**. The instance name must be unique within your Google Cloud project.
4. Enter a permanent **Instance ID** to identify your instance. The instance ID must be unique within your Google Cloud project.
5. Select **Choose a configuration**, click **Regional** or **Multi-region**. You can also compare the configuration details, for understanding the pricing.
6. Optionally – to add a read-only replica, do the following
 - a. Click on Read-Only Replicas
 - b. Select Add Read-Only Replica
 - c. Select a region and the number accordingly
7. Under the Configure Computing Capacity, choose one of the below –
 - a. Processing Units (or)
 - b. Number of Nodes
8. Choose a scaling methodology – Manual Allocation (or) Autoscaling
9. Click on Create, to create the Spanner Instance
10. Trigger the below command on the Google Cloud Terminal to list the spanner instances –
 - a. **gcloud spanner instances list**
11. You can edit the instance configuration details either through the console, or by triggering the gcloud commands. For example, if you want to edit the number of nodes, use the below command –
 - a. **gcloud spanner instances update “your-instance-id --nodes=”desired-node-count”**